

# day\_20\_03Assignment

January 22, 2023

Subject of notebook : Apply machine learning algorithm on any dataset other than titanic or iris

Name of the author : Qadir Shahbaz

Where to contact : qadir\_shahbaz@yahoo.co.uk

date : 22/01/2023

## 0.1 import Libraries

```
[ ]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from sklearn.model_selection import train_test_split
```

## 0.2 Import dataset

```
[ ]: df = sns.load_dataset("penguins")
```

## 0.3 Understand the data

```
[ ]: df.head()
```

```
[ ]: species      island  bill_length_mm  bill_depth_mm  flipper_length_mm  \
0  Adelie  Torgersen      39.1           18.7           181.0
1  Adelie  Torgersen      39.5           17.4           186.0
2  Adelie  Torgersen      40.3           18.0           195.0
3  Adelie  Torgersen      NaN           NaN           NaN
4  Adelie  Torgersen      36.7           19.3           193.0

      body_mass_g  sex
0         3750.0  Male
1         3800.0 Female
2         3250.0 Female
3            NaN   NaN
4         3450.0 Female
```

```
[ ]: df.columns
```

```
[ ]: Index(['species', 'island', 'bill_length_mm', 'bill_depth_mm',
          'flipper_length_mm', 'body_mass_g', 'sex'],
          dtype='object')
```

```
[ ]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   species               344 non-null   object
1   island                344 non-null   object
2   bill_length_mm        342 non-null   float64
3   bill_depth_mm         342 non-null   float64
4   flipper_length_mm     342 non-null   float64
5   body_mass_g           342 non-null   float64
6   sex                   333 non-null   object
dtypes: float64(4), object(3)
memory usage: 18.9+ KB
```

## 0.4 Machine learning algorithms

```
[ ]: X = df[["bill_length_mm", "bill_depth_mm", "flipper_length_mm" ]]
X.fillna(value= X.mean(), inplace = True)
y = df["species"]
```

```
C:\Users\qadir\AppData\Local\Temp\ipykernel_13016\467528285.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X.fillna(value= X.mean(), inplace = True)
```

```
C:\Users\qadir\AppData\Local\Temp\ipykernel_13016\467528285.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```
X.fillna(value= X.mean(), inplace = True)
```

```
C:\Users\qadir\AppData\Local\Temp\ipykernel_13016\467528285.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)  
X.fillna(value= X.mean(), inplace = True)

```
[ ]: X.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 344 entries, 0 to 343
Data columns (total 3 columns):
#   Column                Non-Null Count  Dtype
---  -
0   bill_length_mm        344 non-null   float64
1   bill_depth_mm         344 non-null   float64
2   flipper_length_mm     344 non-null   float64
dtypes: float64(3)
memory usage: 8.2 KB
```

```
[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, precision_score,
    recall_score
from sklearn.model_selection import train_test_split
```

```
[ ]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    random_state=42)

models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    'KNN']
```

```
[ ]: models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    accuracy = accuracy_score(y_test, y_pred)
    models_scores.append([model_name, accuracy])
```

```
c:\Users\qadir\AppData\Local\Programs\Python\Python310\lib\site-
packages\sklearn\linear_model\_logistic.py:458: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:

```

https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
n_iter_i = _check_optimize_result(

```

```

[ ]: sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)
for model in sorted_models:
    print("Accuracy Score: ",f'{model[0]} : {model[1]:.2f}')

```

```

Accuracy Score: Logistic Regression : 0.97
Accuracy Score: Random Forest : 0.97
Accuracy Score: Decision Tree : 0.93
Accuracy Score: KNN : 0.93
Accuracy Score: SVM : 0.77

```

```

[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    ↪RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    ↪'KNN']
models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    Precision = precision_score(y_test,y_pred, average='micro')
    models_scores.append([model_name,Precision])

```

```

[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    ↪RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    ↪'KNN']
models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    Recall = recall_score(y_test, y_pred)
    models_scores.append([model_name,Recall])

```

```

[ ]: models = [LogisticRegression(), SVC(), DecisionTreeClassifier(),
    ↪RandomForestClassifier(), KNeighborsClassifier()]
model_names = ['Logistic Regression', 'SVM', 'Decision Tree', 'Random Forest',
    ↪'KNN']
models_scores = []
for model, model_name in zip(models, model_names):
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)
    F1 = f1_score(y_test, y_pred)
    models_scores.append([model_name,F1])

sorted_models = sorted(models_scores, key=lambda x: x[1], reverse=True)

```

```
for model in sorted_models:  
    print("F1 Score: ",f'{model[0]} : {model[1]:.2f}')
```